

# Algorithm Engineering

Md. Atiqur Rahman  
Segment Tree

# Managing Electricity Usage in a Smart City

A smart city tracks the electricity usage (in kilowatt-hours) of  $N$  households along a single street. The city's control center needs to process two types of operations efficiently:

- **Update Consumption:** Sometimes a household's meter reading is corrected after an inspection. This means replacing the recorded consumption for that household with a new value.
- **Neighborhood Consumption Report:** City officials may request the total electricity usage for a continuous block of houses (for example, houses 10 through 25) to analyze neighborhood consumption patterns.

Given the number of houses and their initial recorded consumption values, you must process  $Q$  operations and provide answers for all neighborhood consumption reports.

# Input Format

N Q

C[1] C[2] ... C[N]

query\_1

query\_2

...

query\_Q

## A test case

5 5

2 1 3 4 5

2 1 5

1 3 10

2 2 4

1 5 7

2 4 5

Each query is in one of the following forms:

- 1 H V  $\rightarrow$  Update household H's consumption to V kWh.
- 2 L R  $\rightarrow$  Output the total consumption for all households from L to R (inclusive).

# Skeleton Algorithm

We need to maintain the following query properties,

- $\text{update}(H, V)$ : set consumption of house  $H$  to  $V$ .
- $\text{query}(L, R)$ : sum consumption from house  $L$  to  $R$  (inclusive).
- Both update and query definition must be implemented in a way so that it requires  $O(\log N)$ .

The layout to save the data,

- Use an iterative segment tree stored in an array of size  $2 \cdot P$ , where  $P$  is the next power of two  $\geq N$ .
- Leaves (original array values) live at indices  $[P .. P+N-1]$ .
- Parents store the sum of their two children.

Building the aforementioned tree,

- Copy  $C[i]$  into  $\text{tree}[P+i]$ .
- For  $i = P-1$  down to  $1$ :  $\text{tree}[i] = \text{tree}[2 \cdot i] + \text{tree}[2 \cdot i + 1]$ .

# Pseudo Code To Build The Tree

```
def next_pow2(n):
```

```
    p = 1
```

```
    while p < n:
```

```
        p <<= 1
```

```
    return p
```

```
P = next_pow2(N)
```

```
tree = [0] * (2 * P)
```

```
for i in range(N):
```

```
    tree[P + i] = C[i]
```

```
for i in range(P - 1, 0, -1):
```

```
    tree[i] = tree[2 * i] + tree[2 * i + 1]
```

# Continued

Point update [ $O(\log n)$  complexity],

- Convert house index  $H$  (1-based) to leaf position  $pos = P + (H-1)$ .
- Set  $tree[pos] = V$ .
- Climb up:  $pos \mathrel{/=} 2$  each step, recomputing  $tree[pos] = tree[2*pos] + tree[2*pos+1]$ .

Range Sum Query [ $O(\log n)$  complexity],

- Convert  $[L, R]$  (1-based) to 0-based half-open in the leaf layer:  $l = P + (L-1)$ ,  $r = P + R$  (note:  $r$  is exclusive in this iterative style).
- While  $l < r$ :
  - If  $l$  is a right child, include  $tree[l]$  and  $l += 1$ .
  - If  $r$  is a right boundary, move left:  $r -= 1$  and include  $tree[r]$ .
  - Move both up:  $l \mathrel{/=} 2$ ,  $r \mathrel{/=} 2$ .
- The accumulated sum is the answer.

# Pseudo Code For Update

```
def update(h, v):
```

```
    pos = P + (h - 1)
```

```
    tree[pos] = v
```

```
    pos //= 2
```

```
    while pos >= 1:
```

```
        tree[pos] = tree[2 * pos] + tree[2 * pos + 1]
```

```
        pos //= 2
```

# Pseudo Code For Update

```
def query(l, r):
```

```
    l = P + (l - 1)
```

```
    r = P + r
```

```
    res = 0
```

```
    while l < r:
```

```
        if l & 1:
```

```
            res += tree[l]
```

```
            l += 1
```

```
        if r & 1:
```

```
            r -= 1
```

```
            res += tree[r]
```

```
        l //= 2
```

```
        r //= 2
```

```
    return res
```



# Some questions

## 1. Why N is the “next power of two” (P)?

- In a complete binary tree, all leaves are at the same depth.
- If N is not a power of two, the tree won't be perfectly balanced; some leaves will be missing.
- Padding up to the next power of two (P) makes indexing predictable — the leaves occupy exactly positions P through P+P-1.
- This also allows us to use simple parent/child index math:
  - Parent of node  $i \rightarrow i // 2$
  - Left child of  $i \rightarrow 2 * i$
  - Right child of  $i \rightarrow 2 * i + 1$

# Some questions

## 2. Why $2 * P$ is the total size?

- The tree has:
  - $P$  leaves (actual data + padding)
  - $P - 1$  internal nodes (sums of children)
  - Total nodes =  $P + (P - 1) = 2 * P - 1$
- We allocate  $2 * P$  slots instead of  $2 * P - 1$  just to make indexing clean and 1-based:
  - Node 1 = root
  - Leaves start at index  $P$
  - We ignore index 0 for simplicity.

# Continued

What is the total number of nodes (including leaves), if a perfect binary tree has  $P$  leaves node?


Always  $2P-1$ .

Why  $2P$  memory allocation works?

In the extreme case scenario, the next power of 2 of  $n$  can be  $2^{\lceil \log_2 n \rceil}$ . So our  $P$  can be  $2^{\lceil \log_2 n \rceil}$  that is equal to  $n$ , and if we allocate  $2P$  that has  $P$  leaves. Since,  $2P-1 < 2P$ .

Two properties of perfect binary tree:

1. Every internal node has exactly two children.
2. All leaf nodes are at the same level (or depth).



# Some questions

## 3. Why make $r$ exclusive?

We want to query  $[L, R]$  inclusive from the user's perspective, but the iterative algorithm internally works on  $[l, r)$  — half-open intervals — because:

- It simplifies the loop: `while l < r`
- It lets us stop exactly when the two pointers meet without double-counting.
- It matches how array slicing works in Python (`[start, end)`).

# Continued

## 4. Significance of $l$ being a right child ( $l \& 1$ )

In the tree array:

- Left child indices are even.
- Right child indices are odd.

If  $l$  is a right child, it means:

- Its parent's segment starts before  $l$ .
- So, if we included the whole parent, we'd be including stuff before the query range.
- Therefore, we take this single node (add it to the sum), and then move  $l$  to the next segment ( $l += 1$ ).

# Conituned

## 5. Significance of $r$ being at a right boundary ( $r \& 1$ )

Remember  $r$  is exclusive — it points to the segment after the range.

If  $r$  is a right child, that means:

- The segment just before  $r$  is the last part of the query.
- So we first move back ( $r -= 1$ ), then take that node's value.

# Recursive Segment Tree


To implement the right side segment tree we need  $4n$  allocation of memory space.

Now, tell me, is  $2P > 4n$ ?

Let,  $2^{\{x\}} \leq n < 2^{\{x+1\}}$ , now highest value of  $n$  can be  $2^{\{x+1\}} - 1$ , in that case  $P$  will be  $2^{\{x+1\}}$ . So,  $2P$  will be  $2^{\{x+2\}}$ .

Similarly, 4 which is  $2^{\{2\}}$ , multiplied to  $2^{\{x+1\}} - 1$ , we get  $2^{\{x+3\}} - 2^{\{2\}}$ ,

These two terms will only be equal if  $x$  is 0 other than that the second term will always be higher. Then, why recursive approach requires higher memory?



# Build Function

If  $tl = 0$ ,  $tr = n-1$ ,  $v = 1$ ,  $n = 5$  and  $a = [1, 3, -2, 8, -7]$ ,

Write all the elements reside inside  $t$  array after executing the right side function.

```
int t[4*n] = {-1};
```

```
void build(int a[], int v, int tl, int tr) {  
    if (tl == tr) {  
        t[v] = a[tl];  
    } else {  
        int tm = (tl + tr) / 2;  
        build(a, v*2, tl, tm);  
        build(a, v*2+1, tm+1, tr);  
        t[v] = t[v*2] + t[v*2+1];  
    }  
}
```



# Is $4n$ enough?

What is the relation between height of a perfect binary tree with the number of nodes it has?

Let,  $h$  be the height of the the perfect binary tree then, number of nodes is  $2^h - 1$ .

What is the length of the highest range we can find sum in a segment tree?  
 $n$ .

Now, this range of size  $n$  is divided by 2 for each level of the segment tree. Thus, the maximum height of the tree can be  $\text{upper\_bound}(\log_2 n) + 1$ . So even if we consider our segment tree as perfect, we have  $2^{\text{upper\_bound}(\log_2 n) + 1} - 1$  nodes.

$$2^{\text{upper\_bound}(\log_2 n) + 1} - 1 < 2^{\text{upper\_bound}(\log_2 n) + 1} < 2 \times 2^{\text{upper\_bound}(\log_2 n) + 1}$$

$$\Rightarrow 2 \times 2^{\text{upper\_bound}(\log_2 n)} \times 2 = 4 \times n$$

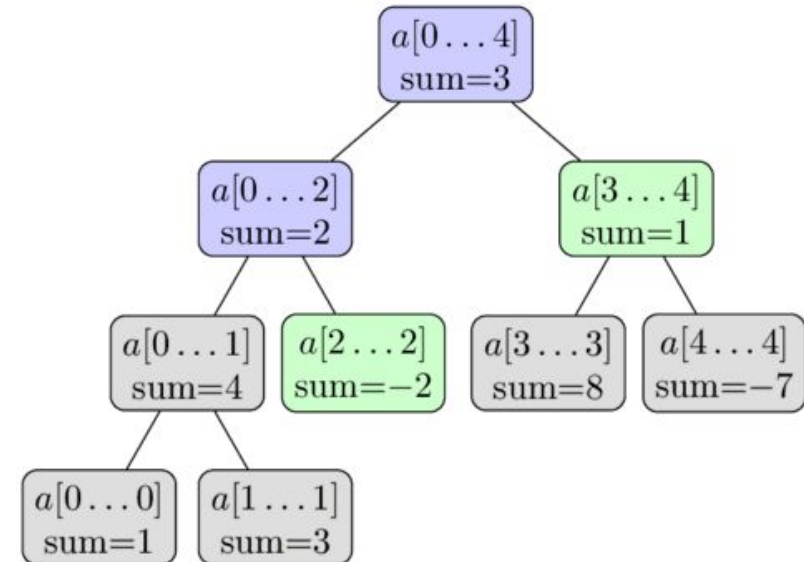
Since,  $4n$  is higher than the maximum nodes possible, it is enough.

# Sum Queries

$v=1$ ,  $tl=0$ ,  $tr=4$ ,  $l=2$ ,  $r=4$ ,

Which nodes in the segment tree, I will visit if we have the above values for the right side sum function?


```
int sum(int v, int tl, int tr, int l, int r) {  
    if (l > r)  
        return 0;  
    if (l == tl && r == tr) {  
        return t[v];  
    }  
    int tm = (tl + tr) / 2;  
    return sum(v*2, tl, tm, l, min(r, tm))  
        + sum(v*2+1, tm+1, tr, max(l, tm+1), r);  
}
```



# Update Queries

Which nodes of the segment tree, I will visit if  $v=1$ ,  $tl=0$ ,  $tr=4$ ,  $pos = 2$ ,  $new\_val = 3$ ?

```
void update(int v, int tl, int tr, int pos, int new_val) {  
    if (tl == tr) {  
        t[v] = new_val;  
    } else {  
        int tm = (tl + tr) / 2;  
        if (pos <= tm)  
            update(v*2, tl, tm, pos, new_val);  
        else  
            update(v*2+1, tm+1, tr, pos, new_val);  
        t[v] = t[v*2] + t[v*2+1];  
    }  
}
```



# Memory efficient Implementation

If you look at the array  $t$  you can see that it follows the numbering of the tree nodes in the order of a BFS traversal (level-order traversal). Using this traversal the children of vertex  $v$  are  $2v$  and  $2v + 1$  respectively.

However if  $n$  is not a power of two, this method will skip some indices and leave some parts of the array  $t$  unused. The memory consumption is limited by  $4n$ , even though a Segment Tree of an array of  $n$  elements requires only  $2n - 1$  vertices.

However it can be reduced. We renumber the vertices of the tree in the order of a pre-order traversal, and we write all these vertices next to each other.

If a node is responsible for  $[l, r]$ , the number of vertices it needs is  $2(r-l+1) - 1$ .

Now, left child's node is  $v+1$ , but the number of vertices it needs is,  $2(\text{mid}-l+1)-1$ , so the right child's vertex number  $v+2(\text{mid}-l+1)$ .

**Implement the build, query and update function using the approach explained in this slide.**